

Automated High-Level Loop Fusion for FLAME Algorithms

Krzysztof A. Drewniak

Carnegie Mellon University

June TODO, 2018

High-level loop fusion

- ▶ Problems often are a series of subproblems
- ▶ Combining subalgorithms often helps performance
- ▶ Goal: find all the fused algorithms for a problem
- ▶ Compilers know too many details - need a high level approach

FLAME algorithms, loop invariants

- ▶ FLAME = Formal Linear Algebra Methods Eenvironments
- ▶ Provably correct algorithms from spec
- ▶ Algorithms \Leftrightarrow loop invariants
- ▶ We know how to:
 - ▶ Autogenerate algorithm/code from loop invariant
 - ▶ Autogenerate all possible loop invariants
 - ▶ Identify when fusion is possible (in theory)

What we add

- ▶ Autogenerate all sets of fusible loop invariants
- ▶ Input is *partitioned matrix expression* — indicates needed computations
- ▶ Can be used to generate code

Section 2

FLAME

Goal

Want to compute

$$\tilde{A} = \mathcal{F}(\hat{A}, \underbrace{\dots}_O)$$

\hat{A} and \tilde{A} share memory (A).

Initially, $A = \hat{A}$.

At termination, $A = \tilde{A}$.

$$\tilde{A} = CHOL(\hat{A})$$

Algorithm structure

partition $A \rightarrow \left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right)$

where $\dim(A_{TL}) = 0 \times 0$

do until $\dim(A_{TL}) = n \times n$

repartition $\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} A_{00} & a_{01} & A_{02} \\ \hline a_{10}^T & \alpha_{11} & a_{12}^T \\ \hline A_{02} & a_{21} & A_{22} \end{array} \right)$

\vdots] loop body

continue with $\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c|c|c} A_{00} & a_{01} & A_{02} \\ \hline a_{10}^T & \alpha_{11} & a_{12}^T \\ \hline A_{20} & a_{21} & A_{22} \end{array} \right)$

enddo

Algorithm example

partition $A \rightarrow \left(\begin{array}{c|c} A_{TL} & * \\ \hline A_{BL} & A_{BR} \end{array} \right)$

where $\dim(A_{TL}) = 0 \times 0$

do until $\dim(A_{TL}) = n \times n$

repartition $\left(\begin{array}{c|c} A_{TL} & * \\ \hline A_{BL} & A_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} A_{00} & * & * \\ \hline a_{10}^T & \alpha_{11} & * \\ \hline A_{02} & a_{21} & A_{22} \end{array} \right)$

$$\alpha_{11} := \sqrt{\alpha_{11}}$$

$$a_{21} := a_{21} / \alpha_{11}$$

$$A_{22} := A_{22} - a_{21} a_{21}^T$$

continue with $\left(\begin{array}{c|c} A_{TL} & * \\ \hline A_{BL} & A_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c|c|c} A_{00} & * & * \\ \hline a_{10}^T & \alpha_{11} & * \\ \hline A_{20} & a_{21} & A_{22} \end{array} \right)$

enddo

Partitioned Matrix Expressions

- ▶ Take A (and maybe other stuff), split it into regions.
- ▶ Lines between regions move during algorithm

$$\left(\frac{\tilde{A}_{TL} = \mathcal{F}_{TL}(\hat{A}, \dots) \parallel \tilde{A}_{TR} = \mathcal{F}_{TR}(\hat{A}, \dots)}{\tilde{A}_{BL} = \mathcal{F}_{BL}(\hat{A}, \dots) \parallel \tilde{A}_{BR} = \mathcal{F}_{BR}(\hat{A}, \dots)} \right)$$

$$\left(\frac{\tilde{A}_{TL} = \text{CHOL}(\hat{A}_{TL}) \parallel \quad \quad \quad *}{\tilde{A}_{BL} = \hat{A}_{BL} \tilde{A}_{TL}^{-T} \parallel \tilde{A}_{BR} = \text{CHOL}(\hat{A}_{BR} - \tilde{A}_{BL} \tilde{A}_{BL}^T)} \right)$$

Loop invariants

- ▶ Find f_R and f'_R so $\mathcal{F}_R(\hat{A}) = f'(f(\hat{A}))$.
- ▶ f_R is loop invariant for R , f'_R is remainder
- ▶ Invariant for algorithm is an invariant per region
- ▶ Completely determine algorithm

This is a loop invariant

Starting from Cholesky's PME:

$$\left(\frac{\tilde{A}_{TL} = CHOL(\hat{A}_{TL})}{\tilde{A}_{BL} = \hat{A}_{BL} \tilde{A}_{TL}^{-T}} \parallel \frac{*}{\tilde{A}_{BR} = CHOL(\hat{A}_{BR} - \tilde{A}_{BL} \tilde{A}_{BL}^T)} \right)$$

We obtain

$$\left(\frac{A_{TL} = CHOL(\hat{A}_{TL})}{A_{BL} = \hat{A}_{BL} \tilde{A}_{TL}^{-T}} \parallel \frac{*}{A_{BR} = \hat{A}_{BR} - \tilde{A}_{BL} \tilde{A}_{BL}^T} \right)$$

As are these

$$\left(\begin{array}{c|c} A_{TL} = CHOL(\hat{A}_{TL}) & * \\ \hline A_{BL} = \hat{A}_{BL} \tilde{A}_{TL}^{-T} & A_{BR} = \hat{A}_{BR} \end{array} \right)$$

$$\left(\begin{array}{c|c} A_{TL} = CHOL(\hat{A}_{TL}) & * \\ \hline A_{BL} = \hat{A}_{BL} & A_{BR} = \hat{A}_{BR} \end{array} \right)$$

But not these

$$\left(\frac{A_{TL} = CHOL(\hat{A}_{TL}) \parallel *}{A_{BL} = \hat{A}_{BL} \tilde{A}_{TL}^{-T} \parallel A_{BR} = CHOL(\hat{A}_{BR} - \tilde{A}_{BL} \tilde{A}_{BL}^T)} \right)$$

$$\left(\frac{A_{TL} = \hat{A}_{TL} \parallel *}{A_{BL} = \hat{A}_{BL} \parallel A_{BR} = \hat{A}_{BR}} \right)$$

Or this

$$\left(\begin{array}{c|c} A_{TL} = \hat{A}_{TL} & * \\ \hline A_{BL} = \hat{A}_{BL} \tilde{A}_{TL}^{-T} & A_{BR} = \hat{A}_{BR} \end{array} \right)$$

Tasks

- We need to specify split points

$$\left(\begin{array}{c|c} \tilde{A}_{TL} :=_O CHOL(\hat{A}_{TL}) & * \\ \hline \tilde{A}_{BL} := \hat{A}_{BL} \tilde{A}_{TL}^{-T} & \begin{array}{l} A_{BR,0} := \hat{A}_{BR} - \tilde{A}_{BL} \tilde{A}_{BL}^T; \\ \tilde{A}_{BR} :=_O CHOL(A_{BR,0}) \end{array} \end{array} \right)$$

$$\left(\begin{array}{c|c} \tilde{A}_{TL} :=_O \hat{A}_{TL}^{-1} & * \\ \hline \begin{array}{l} A_{BL,(0,a)} := (\hat{A}_{BL} \vee A_{BL,(0,b)}) \cdot \tilde{A}_{TL}; \\ A_{BL,(0,b)} := -\hat{A}_{BR}^{-1} \cdot (\hat{A}_{BL} \vee A_{BL,(0,a)}) \end{array} & \tilde{A}_{BR} :=_O \hat{A}_{BR}^{-1} \end{array} \right)$$

More abstractly

The code translates tasks to

$$\left(\frac{A_{TL,\top} :=_O \{A_{TL,\perp}\}}{A_{BL,\top} := \{A_{BL,\perp}, A_{TL,\top}\}} \parallel \frac{*}{A_{BR,0} := \{A_{BR,\perp} A_{BL,\top}; A_{BR,\top} :=_O \{A_{BR,0}\}} \right)$$

$$\left(\frac{A_{TL,\top} :=_O \{A_{TL,\perp}\}}{A_{BL,(0,a)} := \{A_{BL,\perp} \vee A_{BL,(0,b)}, A_{TL,\top}\}; A_{BL,(0,b)} := \{A_{BR,\perp}, A_{BL,\perp} \vee A_{BL,(0,a)}\}} \parallel \frac{*}{A_{BR,\top} :=_O \{A_{BR,\perp}\}} \right)$$

Dependencies, v2

- ▶ $A_{R,\sigma}$ is before $A_{R',\sigma'}$ if:
 - ▶ $R \neq R'$ (different regions) or
 - ▶ $\sigma = \perp$ and $\sigma' \neq \perp$
 - ▶ $\sigma \neq \top$ and $\sigma' = \top$
 - ▶ $\sigma = m$ (or (m, x)) and $\sigma' = n$ (or (n, y)), $m < n$
 - ▶ $\sigma = (n, x)$ and $\sigma' = (n, y)$, and $x \neq y$
- ▶ If anything from am or is before, all of it is
- ▶ Invariant/remainder split valid if:
 - ▶ All past task inputs before all future task outputs
 - ▶ All past task outputs not after all future task inputs

Finding all invariants

1. Pick a past/future split for each region
2. Check if the loop can make progress
3. Check for dependency validity

Section 3

Loop fusion

States of a region

Fully computed All tasks in the invariant

Uncomputed All tasks in the remainder

Partially computed Everything else

The fusion problem

$$\left. \begin{array}{lcl} \tilde{A}^0 & = & \mathcal{F}^0(\hat{A}^0) \\ \tilde{A}^1 & = & \mathcal{F}^1(\hat{A}^1) \\ & \vdots & \\ \tilde{A}^{n-1} & = & \mathcal{F}^{n-1}(\hat{A}^{n-1}) \end{array} \right\} \tilde{A}^{n-1} = \mathcal{F}(\hat{A}^0)$$

where $\hat{A}^{i+1} = \tilde{A}^i$

Fusion conditions

$$\hat{A}_{\mathbf{R}}^{i+1} = \tilde{A}_{\mathbf{R}}^i \text{ if needed}$$

- ▶ \mathcal{F}^i 's invariant needs $R \Rightarrow \mathcal{F}_R^{j < i}$ fully computed
- ▶ \mathcal{F}^i 's remainder needs $R \Rightarrow \mathcal{F}_R^{j > i}$ uncomputed

An example: Cholesky + lower-triangular solve

Cholesky algorithms.

$$\begin{pmatrix} \frac{L_{TL} = CHOL(\hat{L}_{TL})}{L_{BL} = \hat{L}_{BL}} \parallel * \\ L_{BR} = \hat{L}_{BR} \end{pmatrix}$$

$$\begin{pmatrix} \frac{L_{TL} = CHOL(\hat{L}_{TL})}{L_{BL} = \hat{L}_{BL} \tilde{L}_{TL}^{-T}} \parallel * \\ L_{BR} = \hat{L}_{BR} \end{pmatrix}$$

$$\begin{pmatrix} \frac{L_{TL} = CHOL(\hat{L}_{TL})}{L_{BL} = \hat{L}_{BL} \tilde{L}_{TL}^{-T}} \parallel * \\ L_{BR} = \hat{L}_{BR} - \tilde{L}_{BL} \tilde{L}_{BL}^T \end{pmatrix}$$

Lower triangular solve algorithms.

$$\begin{pmatrix} \frac{B_T = L_{TL} \setminus \hat{B}_T}{B_B = \hat{B}_B} \end{pmatrix}$$

$$\begin{pmatrix} \frac{B_T = L_{TL} \setminus \hat{B}_T}{B_B = \hat{B}_B - L_{BL} \tilde{B}_T} \end{pmatrix}$$

5 fused algorithms. (All combinations fuse except one.)

An example: Cholesky + lower-triangular solve

Cholesky algorithms.

$$\begin{aligned}
 & \left(\begin{array}{c|c} L_{TL} = CHOL(\hat{L}_{TL}) & * \\ \hline L_{BL} = \hat{L}_{BL} & L_{BR} = \hat{L}_{BR} \end{array} \right) \\
 & \left(\begin{array}{c|c} L_{TL} = CHOL(\hat{L}_{TL}) & * \\ \hline L_{BL} = \hat{L}_{BL} \tilde{L}_{TL}^{-T} & L_{BR} = \hat{L}_{BR} \end{array} \right) \\
 & \left(\begin{array}{c|c} L_{TL} = CHOL(\hat{L}_{TL}) & * \\ \hline L_{BL} = \hat{L}_{BL} \tilde{L}_{TL}^{-T} & L_{BR} = \hat{L}_{BR} - \tilde{L}_{BL} \tilde{L}_{BL}^T \end{array} \right)
 \end{aligned}$$

Lower triangular solve algorithm

$$\begin{aligned}
 & \left(\begin{array}{c} B_T = L_{TL} \setminus \hat{B}_T \\ \hline B_B = \hat{B}_B \end{array} \right) \\
 & \left(\begin{array}{c} B_T = L_{TL} \setminus \hat{B}_T \\ \hline B_B = \hat{B}_B - L_{BL} \tilde{B}_T \end{array} \right)
 \end{aligned}$$

5 fused algorithms. (All combinations fuse except one.)

We can't go further

Consider:

$$\begin{aligned}
 L &:= CHOL(L) & \left(\begin{array}{c|c} \tilde{A}_{TL} :=_O CHOL(\hat{A}_{TL}) & * \\ \hline \tilde{A}_{BL} := \hat{A}_{BL} \tilde{A}_{TL}^{-T} & \begin{array}{l} A_{BR,0} := \hat{A}_{BR} - \tilde{A}_{BL} \tilde{A}_{BL}^T; \\ \tilde{A}_{BR} :=_O CHOL(A_{BR,0}) \end{array} \end{array} \right) \\
 T &:= L^{-1}B & \left(\begin{array}{c} \tilde{T}_T :=_O TRSV(\hat{L}_{TL}, B_T) \\ \hline T_{B,0} := \hat{T}_B - \hat{L}_{BL} \tilde{T}_T \\ \tilde{T}_B :=_O TRSV(\hat{L}_{BR}, T_{B,0}) \end{array} \right) \\
 X &:= L^{-T}B & \left(\begin{array}{c} X_{T,0} := \hat{X}_T - \hat{L}_{BL}^T \tilde{X}_B \\ \hline \tilde{X}_T :=_O TRSV(\hat{L}_{BR}, X_{T,0}) \\ \hline X_B := TRSV(\hat{L}_{BR}, \hat{T}_B) \end{array} \right)
 \end{aligned}$$

- ▶ No fused algorithm (we checked)
- ▶ Top to bottom vs. bottom to top

Strips

- ▶ Strip: sequence of region R from each loop
- ▶ Potentially fusable strip has:
 - ▶ Some number of fully computed regions, then
 - ▶ Optionally, one partially computed region, then
 - ▶ Uncomputed regions



but not



Finding fusable loops

- ▶ Search through potentially fusable strips

$$\begin{array}{c} \blacksquare \square \\ \text{Any} \end{array} = \begin{array}{c} \blacksquare \square \\ \text{Any} \end{array}$$



- ▶ Enforce fusion constraints throughout
- ▶ Check all fusable strip-sets to see if each loop has an invariant

Last computed, first uncomputed


- ▶ Track constraints on last computed region C_R (and first uncomputed U_R)
- ▶ Initially, $-1 \leq C_R, U_R \leq n$ (maybe nothing/everything is computed/uncomputed)
- ▶ Past read in loop i : $C_R \geq i - 1$
- ▶ Future read in loop i : $U_R \leq i + 1$
- ▶ When strip is made, set C_R and U_R , add more constraints
- ▶ On failure, backtrack

Multiple matrices

- ▶ Some operations have multiple outputs
- ▶ (Ex. $y = Lx; L = L^{-1}$)
- ▶ All strips must be same length — add empty regions
- ▶ De-dup check from before works

Multiple matrices

- ▶ Last computed or first uncomputed can be followed by empty
- ▶ If so, bound $\{C, U\}_R$ to include the empty regions
- ▶ Needed to make constraints work



Comes from task

- ▶ For things like $LU = A$, tasks write multiple regions
- ▶ To prevent duplicates, use $U_R \leftarrow L_R$ (comes from)
- ▶ If L_R computed, U_R is computed, otherwise not

Another important example

- ▶ Graph problem $C = (AM + (AM)^T) - MM$, where A and M are symmetric
- ▶ $C := (AM + (AM)^T); C := C - MM$ has 56 fused algorithms
- ▶ However, $C = A$ or $C = M$ gives 0 algorithms
 - ▶ Dependencies: $TL \leftrightarrow TR$ and $TR \leftrightarrow BR$
 - ▶ Overwriting one quadrant requires computing everything
 - ▶ **TODO figure**

TODO do an experiment

Conclusions

- ▶ We can automatically find fusable loop invariants
- ▶ This is often helpful
- ▶ This analysis needs to be at this level

Acknowledgments

- ▶ Tze Meng for doing all the theory

Future work

- ▶ Probably not — maybe codegen